

1: Linear regression takes two values, x and y, and procures a line of best fit of y as a function of x with the equation $y = mx + b$. 'm' is the slope and 'b' is the intercept.

Strength: Its very simple to set up and interpret. It works well on liner data. It has a low variance.

Weaknesses: Its simplicity and bias makes it keen to under fit data that doesn't follow a perfectly liner data set.

2: The data used will be "Video Game Sales with Ratings" by Rush Kirubi which can be accessed with the url:

<https://www.kaggle.com/datasets/rush4ratio/video-game-sales-with-ratings> (<https://www.kaggle.com/datasets/rush4ratio/video-game-sales-with-ratings>)

The data frame will be read into the program

```
df <- read.csv("Video_Games_Sales.csv")
```

We will have to remove null data since trying to attach a mean or median value to things like sales and rating don't apply well to video games.

```
df <- na.omit(df)
```

To preview the data, we will peek at it.

```
str(df)
```

```
## 'data.frame': 7017 obs. of 16 variables:
## $ Name : chr "Wii Sports" "Mario Kart Wii" "Wii Sports Resort" "New Super Mario Bros." ...
## $ Platform : chr "Wii" "Wii" "Wii" "DS" ...
## $ Year_of_Release: chr "2006" "2008" "2009" "2006" ...
## $ Genre : chr "Sports" "Racing" "Sports" "Platform" ...
## $ Publisher : chr "Nintendo" "Nintendo" "Nintendo" "Nintendo" ...
## $ NA_Sales : num 41.4 15.7 15.6 11.3 14 ...
## $ EU_Sales : num 28.96 12.76 10.93 9.14 9.18 ...
## $ JP_Sales : num 3.77 3.79 3.28 6.5 2.93 4.7 4.13 3.6 0.24 2.53 ...
## $ Other_Sales : num 8.45 3.29 2.95 2.88 2.84 2.24 1.9 2.15 1.69 1.77 ...
## $ Global_Sales : num 82.5 35.5 32.8 29.8 28.9 ...
## $ Critic_Score : int 76 82 80 89 58 87 91 80 61 80 ...
## $ Critic_Count : int 51 73 73 65 41 80 64 63 45 33 ...
## $ User_Score : chr "8" "8.3" "8" "8.5" ...
## $ User_Count : int 322 709 192 431 129 594 464 146 106 52 ...
## $ Developer : chr "Nintendo" "Nintendo" "Nintendo" "Nintendo" ...
## $ Rating : chr "E" "E" "E" "E" ...
## - attr(*, "na.action")= 'omit' Named int [1:9702] 2 5 6 10 11 13 19 21 22 23 ...
## ... attr(*, "names")= chr [1:9702] "2" "5" "6" "10" ...
```

Some columns need become factors. They are: platform, release year, genre, publisher, rating, developer.

User score needs to be converted into a number.

```
df$Platform <- factor(df$Platform)
df$Year_of_Release <- factor(df$Year_of_Release)
df$Genre <- factor(df$Genre)
df$Publisher <- factor(df$Publisher)
df$Rating <- factor(df$Rating)
df$Developer <- factor(df$Developer)
df$User_Score <- as.numeric(df$User_Score)
```

The user rating will have to be multiplied by 10 due to the strange nature of how Meta critic has Critic scores on a scale from 0-100 but user scores 0-10

```
df$User_Score <- df$User_Score * 10
```

#The Sales will also be multiplied by 100000

```
df$NA_Sales <- df$NA_Sales * 1000000
df$EU_Sales <- df$EU_Sales * 1000000
df$JP_Sales <- df$JP_Sales * 1000000
df$Global_Sales <- df$Global_Sales * 1000000
df$Other_Sales <- df$Other_Sales * 1000000
```

a: The data will be divided into train(80) and test(20)

```
set.seed(123)
i <- sample(1:nrow(df), nrow(df) * .80, replace = FALSE)
train <- df[i,]
test <- df[-i,]
```

b: Various R functions will be used for basic data exploration.

```
names(train)
```

```
## [1] "Name"          "Platform"      "Year_of_Release" "Genre"
## [5] "Publisher"     "NA_Sales"      "EU_Sales"        "JP_Sales"
## [9] "Other_Sales"   "Global_Sales"  "Critic_Score"    "Critic_Count"
## [13] "User_Score"    "User_Count"    "Developer"       "Rating"
```

```
dim(train)
```

```
## [1] 5613 16
```

```
summary(train)
```

```
##      Name      Platform Year_of_Release      Genre
## Length:5613   PS2      : 923   2007      : 483   Action      :1343
## Class :character X360    : 702   2008      : 462   Sports       : 755
## Mode  :character PS3     : 634   2009      : 452   Shooter      : 713
##           PC       : 554   2005      : 437   Role-Playing: 592
##           XB       : 459   2006      : 408   Racing       : 474
##           Wii      : 379   2003      : 401   Platform     : 334
##           (Other):1962 (Other):2970 (Other)     :1402
##           Publisher      NA_Sales      EU_Sales
## Electronic Arts      : 756   Min.      :      0   Min.      :      0
## Ubisoft              : 404   1st Qu.: 60000   1st Qu.: 20000
## Activision           : 388   Median : 150000   Median : 60000
## Sony Computer Entertainment: 257 Mean    : 376574   Mean    : 222147
## THQ                  : 243   3rd Qu.: 380000   3rd Qu.: 200000
## Sega                 : 234   Max.     :41360000 Max.     :28960000
## (Other)              :3331
##      JP_Sales      Other_Sales      Global_Sales      Critic_Score
## Min.      :      0   Min.      :      0   Min.      : 10000   Min.      :13.00
## 1st Qu.:      0   1st Qu.: 10000   1st Qu.: 110000   1st Qu.:62.00
## Median :      0   Median : 20000   Median : 290000   Median :72.00
## Mean     : 62293   Mean     : 77021   Mean     : 738229   Mean     :70.26
## 3rd Qu.: 10000   3rd Qu.: 70000   3rd Qu.: 740000   3rd Qu.:80.00
## Max.     :5320000 Max.     :8450000   Max.     :82530000 Max.     :98.00
##
##      Critic_Count      User_Score      User_Count      Developer
## Min.      : 3.00   Min.      : 5   Min.      : 4.0   EA Sports : 119
## 1st Qu.: 14.00   1st Qu.:65   1st Qu.: 11.0   EA Canada : 113
## Median : 24.00   Median :75   Median : 27.0   Capcom    : 103
## Mean     : 28.75   Mean     :72   Mean     : 173.5   Ubisoft   : 83
## 3rd Qu.: 39.00   3rd Qu.:82   3rd Qu.: 89.0   Konami    : 78
## Max.     :113.00   Max.     :96   Max.     :10665.0 EA Tiburon: 69
##                                     (Other)   :5048
##
##      Rating
## T      :1943
## E      :1683
## M      :1165
## E10+   : 766
##        : 53
## AO     : 1
## (Other): 2
```

```
str(train)
```

```
## 'data.frame':    5613 obs. of  16 variables:
## $ Name          : chr  "Looney Tunes: Back in Action" "Bleach: Soul Resurreccion" "Up" "Dynasty Warriors 4"
## $ Platform      : Factor w/ 17 levels "3DS","DC","DS",...: 8 9 13 8 15 13 11 6 6 3 ...
## $ Year_of_Release: Factor w/ 26 levels "1985","1988",...: 12 20 18 12 17 17 16 16 19 15 ...
## $ Genre         : Factor w/ 12 levels "Action","Adventure",...: 5 3 1 1 9 3 11 9 12 1 ...
## $ Publisher     : Factor w/ 273 levels "10TACLE Studios",...: 264 173 244 239 65 246 66 10 252 10 ...
## $ NA_Sales      : num  250000 270000 220000 630000 80000 340000 600000 0 0 100000 ...
## $ EU_Sales      : num  190000 100000 280000 210000 90000 20000 40000 1120000 270000 0 ...
## $ JP_Sales      : num   0 70000 0 1130000 0 0 0 0 0 0 ...
## $ Other_Sales   : num  60000 50000 60000 130000 20000 30000 70000 30000 50000 10000 ...
## $ Global_Sales  : num  500000 490000 560000 2110000 190000 380000 710000 1150000 310000 100000 ...
## $ Critic_Score  : int   51 58 62 78 52 72 75 92 79 50 ...
## $ Critic_Count  : int    9 34 6 24 35 20 10 40 33 17 ...
## $ User_Score    : num   75 72 80 93 49 84 68 85 52 78 ...
## $ User_Count    : int    6 46 4 201 22 16 8 2360 213 6 ...
## $ Developer     : Factor w/ 1313 levels "", "10tacle Studios, Fusionsphere Systems",...: 1278 909 93 826 868 3
## $ Rating        : Factor w/ 8 levels "", "A0","E","E10+",...: 3 8 4 8 6 8 3 6 4 4 ...
## - attr(*, "na.action")= 'omit' Named int [1:9702] 2 5 6 10 11 13 19 21 22 23 ...
## ... attr(*, "names")= chr [1:9702] "2" "5" "6" "10" ...
```

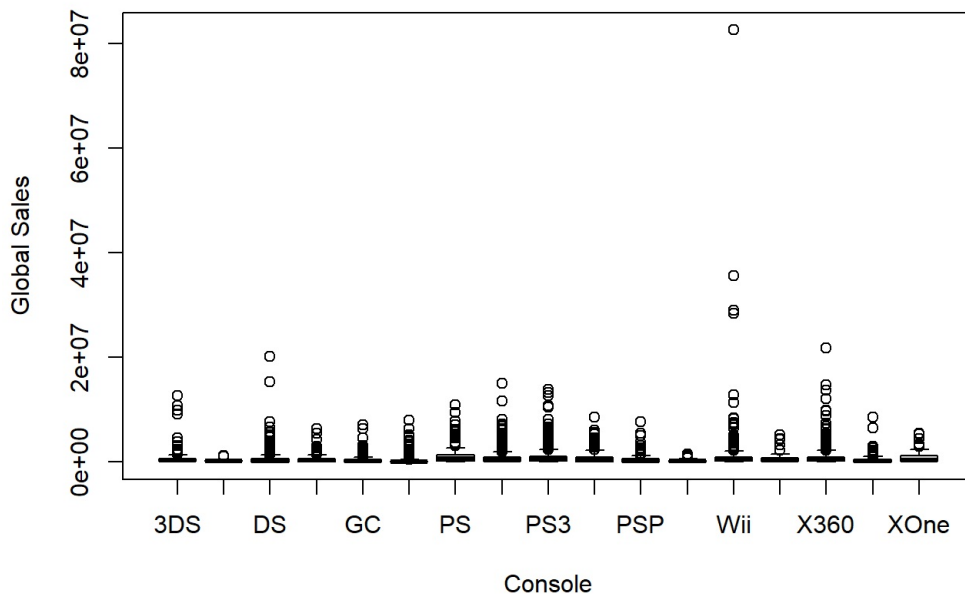
```
head(train)
```

```
##              Name Platform Year_of_Release  Genre
## 3940   Looney Tunes: Back in Action      PS2      2003 Platform
## 4023      Bleach: Soul Resurreccion      PS3      2011 Fighting
## 3549              Up      Wii      2009   Action
## 785      Dynasty Warriors 4      PS2      2003   Action
## 7779      Conflict: Denied Ops     X360      2008  Shooter
## 4978 Naruto: Clash of Ninja Revolution 2    Wii      2008 Fighting
##              Publisher NA_Sales EU_Sales JP_Sales
## 3940 Warner Bros. Interactive Entertainment 250000 190000    0
## 4023      Nippon Ichi Software 270000 100000  70000
## 3549              THQ 220000 280000    0
## 785      Tecmo Koei 630000 210000 1130000
## 7779      Eidos Interactive 80000 90000    0
## 4978      Tomy Corporation 340000 20000    0
## Other_Sales Global_Sales Critic_Score Critic_Count User_Score User_Count
## 3940    60000    500000    51          9        75          6
## 4023    50000    490000    58         34        72         46
## 3549    60000    560000    62          6        80          4
## 785    130000    2110000    78         24        93        201
## 7779    20000    190000    52         35        49         22
## 4978    30000    380000    72         20        84         16
## Developer Rating
## 3940   Warthog      E
## 4023   Racjin      T
## 3549 Asobo Studio  E10+
## 785   Omega Force  T
## 7779 Pivotal Games  M
## 4978   Fighting    T
```

As is noted by the data, there are 5613 entries on video games containing information such as Title, Publisher, Release Year, Sales, ESRB ratings, and Scores.

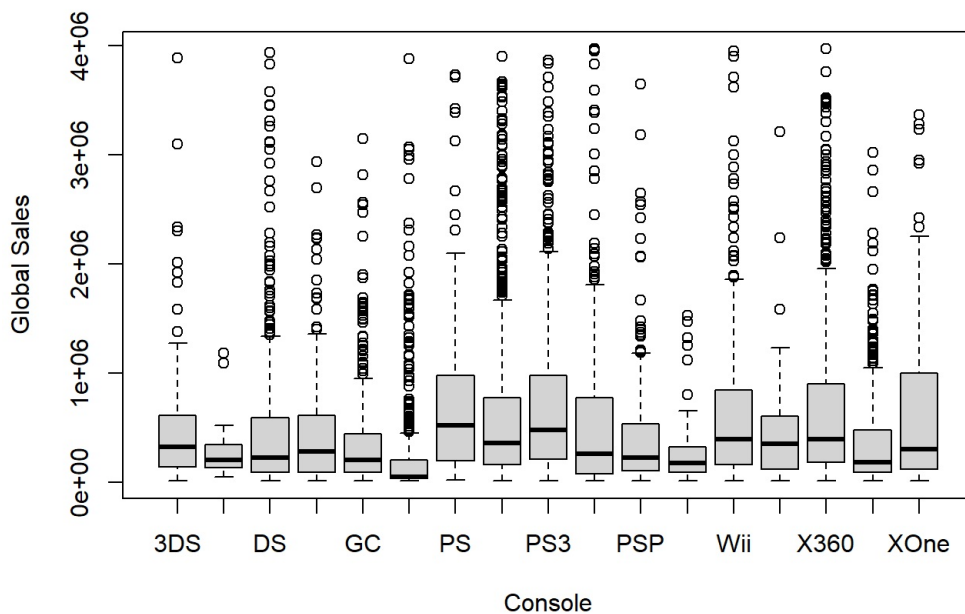
c: Next, some graphs will be used to try and make sense of the data. The first will be Global Sales as a function of the Platform the game released on. To clarify, not all games are platform exclusive and can be released on multiple different platforms.

```
plot(train$Platform, train$Global_Sales, xlab = "Console", ylab = "Global Sales")
```



It would appear there are some extreme outliers that are making the data unusable. As a guess, the cutoff point for realistic entry will be 40 million since only one game sold significantly more than that.

```
train <- subset(train, Global_Sales < 40000000)
plot(train$Platform, train$Global_Sales, xlab = "Console", ylab = "Global Sales")
```



```
levels(train$Platform)
```

```
## [1] "3DS" "DC" "DS" "GBA" "GC" "PC" "PS" "PS2" "PS3" "PS4"
## [11] "PSP" "PSV" "Wii" "WiiU" "X360" "XB" "XOne"
```

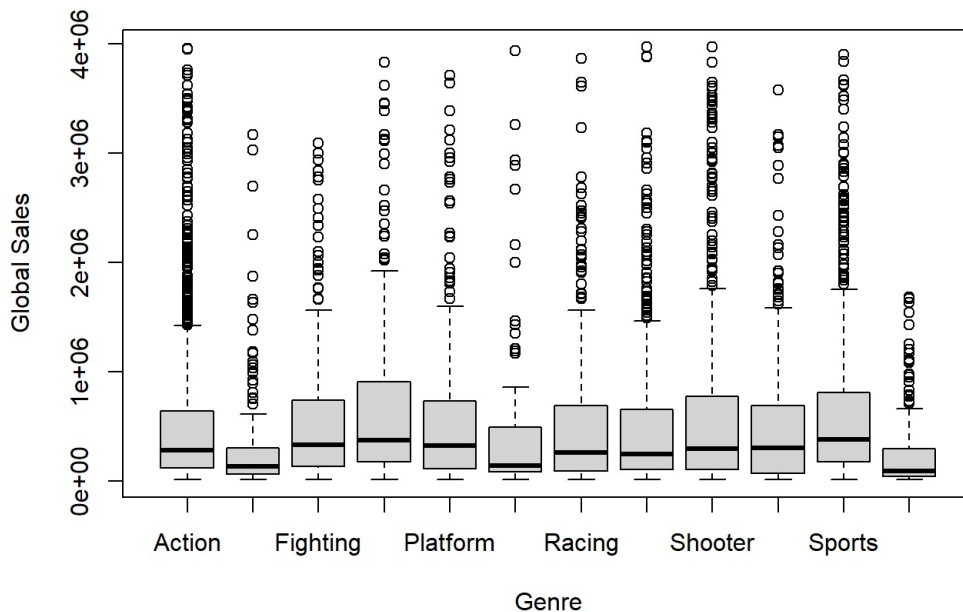
```
WiiUMedian <- (subset(train, Platform == "WiiU"))
median((WiiUMedian$Global_Sales))
```

```
## [1] 350000
```

This data is very rough, but it would appear all the consoles have a rough median of 350 thousand. Some of the larger game companies main colsoles like PlayStation, Xbox, and Nintendo have the highest overall sales. Handheld units such as the PSVita and GameBoy Advanced have less sales.

Next, sales as a function of genre will be checked.

```
plot(train$Genre, train$Global_Sales, xlab = "Genre", ylab = "Global Sales")
```



```
levels(train$Genre)
```

```
## [1] "Action"      "Adventure"    "Fighting"     "Misc"         "Platform"
## [6] "Puzzle"      "Racing"       "Role-Playing" "Shooter"      "Simulation"
## [11] "Sports"      "Strategy"
```

```
ActionMedian <- (subset(train, Genre == "Action"))
median((ActionMedian$Global_Sales))
```

```
## [1] 280000
```

The median is roughly consistent across genres at about 280 thousand. Genres such as Sport and surprisingly Misc sold the most. Strategy and surprisingly Adventure game sold the least.

Next, the cor functions will be used to find trends in the numerical data.

```
cor(train[6:14], use="complete")
```

```
##      NA_Sales  EU_Sales  JP_Sales  Other_Sales  Global_Sales
## NA_Sales    1.0000000  0.57952717  0.13927812  0.53357266  0.8949817
## EU_Sales    0.5795272  1.00000000  0.14577893  0.61162404  0.8265761
## JP_Sales    0.1392781  0.14577893  1.00000000  0.13938471  0.3566772
## Other_Sales  0.5335727  0.61162404  0.13938471  1.00000000  0.7160424
## Global_Sales 0.8949817  0.82657611  0.35667721  0.71604239  1.0000000
## Critic_Score 0.2984213  0.27508787  0.13214737  0.21343413  0.3319263
## Critic_Count 0.3081760  0.32196489  0.14806516  0.22715531  0.3601730
## User_Score   0.1140989  0.07773282  0.14893739  0.05624837  0.1324579
## User_Count   0.1397130  0.30167729  0.02363287  0.17707858  0.2241563
##
##      Critic_Score Critic_Count User_Score User_Count
## NA_Sales    0.2984213    0.3081760  0.11409894  0.13971300
## EU_Sales    0.2750879    0.3219649  0.07773282  0.30167729
## JP_Sales    0.1321474    0.1480652  0.14893739  0.02363287
## Other_Sales 0.2134341    0.2271553  0.05624837  0.17707858
## Global_Sales 0.3319263    0.3601730  0.13245788  0.22415635
## Critic_Score 1.0000000    0.3679886  0.58280018  0.23261475
## Critic_Count 0.3679886    1.0000000  0.18886315  0.28350284
## User_Score   0.5828002    0.1888631  1.00000000  0.02333167
## User_Count   0.2326147    0.2835028  0.02333167  1.00000000
```

NA and EU game sales have a decent cor of .57. This is most likely due to the cultures of the two being similar and thus having similar taste in video games.

Critic score and user score seem to have a decent correlation of .58. This is most likely due to critics grading without bias and knowing more about what games ought to provide. A casual reviewer will be more influenced by their lack of knowledge and just how they felt playing the game.

Unfortunately, it doesn't appear as if user and critic scores have much influence on sales. Critics scores only have a .33 cor and user only have a .13 cor.

d: Unfortunately, it doesn't seem as there are many useful relations with the data set. Even though NA sales and global sales are strongly correlated, the issue is NA sales are being recorded directly within global sales. Critic and user reviews also occur independently of each other, so it's not realistic to put either as a function of the other.

Even though the correlation is poor, sales and a function of critic reviews will be tried as its the only logical connection.

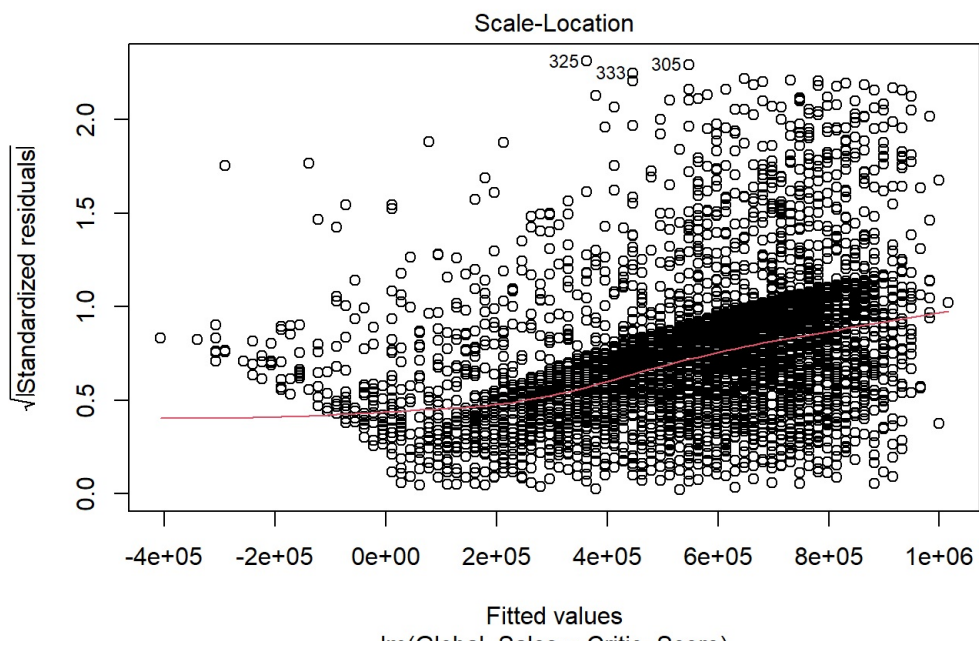
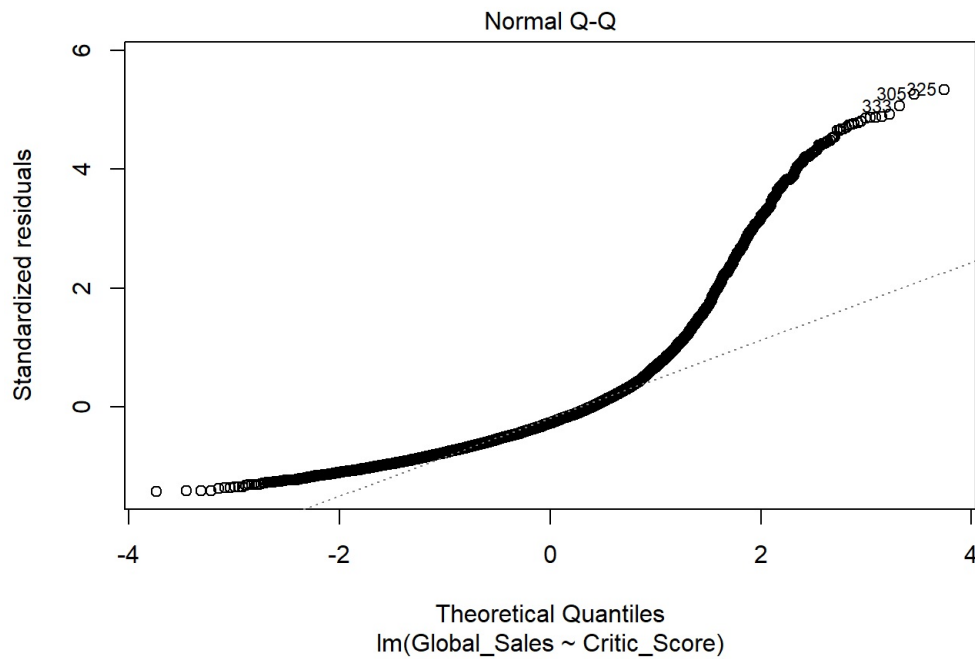
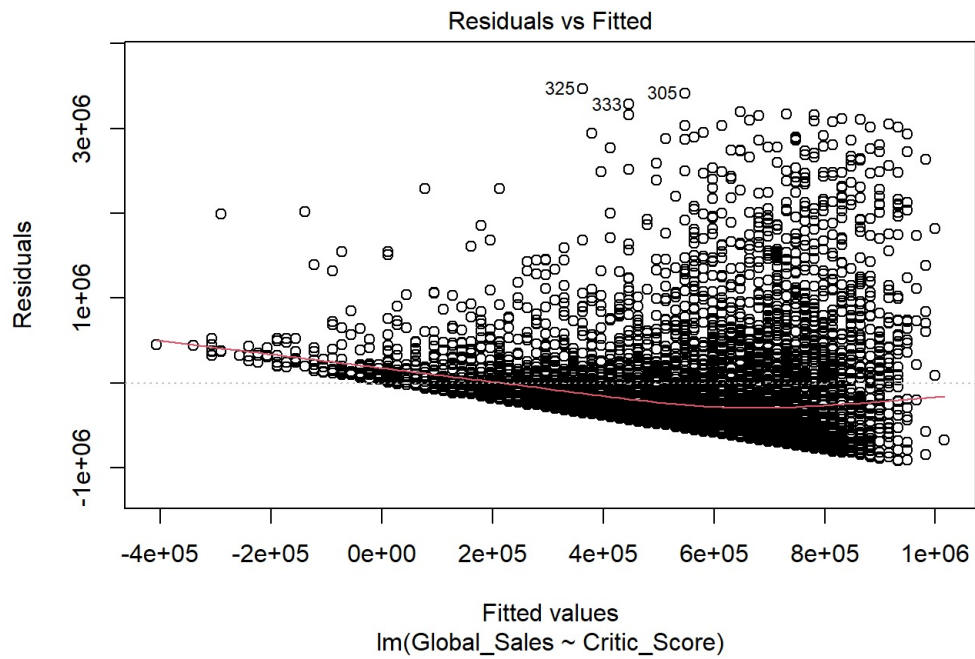
```
lm1 <- lm(Global_Sales~Critic_Score, data=train)
summary(lm1)
```

```
##
## Call:
## lm(formula = Global_Sales ~ Critic_Score, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -922200 -406753 -173745  167596 3467231
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -625359.5    45837.5  -13.64  <2e-16 ***
## Critic_Score   16747.9      644.3    25.99  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 649300 on 5457 degrees of freedom
## Multiple R-squared:  0.1102, Adjusted R-squared:  0.11
## F-statistic: 675.7 on 1 and 5457 DF,  p-value: < 2.2e-16
```

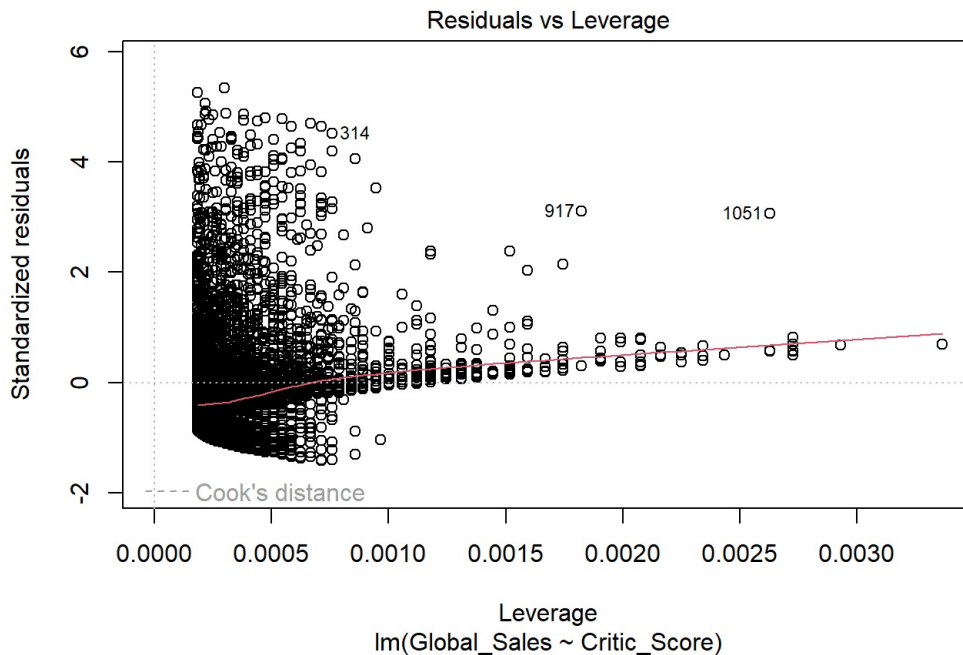
As expected, the min and max from the data vary by several million, and Q1 and Q3 vary by hundreds of thousands. The model suggests a roughly 16 thousand sales increase per critic point awarded with an error range of about half a million. The R squared accuracy is only .11, which is awful. At least the p value is low while the F statistic is high.

e: Even though we know the data is awful, a residual model will be used next.

```
plot(lm1)
```



lm(Global_Sales ~ Critic_Score)



In residual vs. fitted, it doesn't seem

very useful. The red line follows no distinct pattern and The spread of residuals above and below the line aren't even.

In normal Q-Q, as expected, the residuals wildly fly off the line at the 1 mark. This means the residuals aren't normally distributed. It makes sense, considering the abundant amount of outliers in the sales above 1 million.

In scale location, the line starts off decently horizontal, but begins to grow rapidly. Also, the spread is not great. It starts off almost exclusively above the line, but then is mostly below the line.

In residuals vs. leverage, the graph has a lot of influential points since half the data cuts through one of cook's lines.

As expected, the 4 graphs denote a terrible model.

f: Next, the platform will be included to develop a multiple liner regression model.

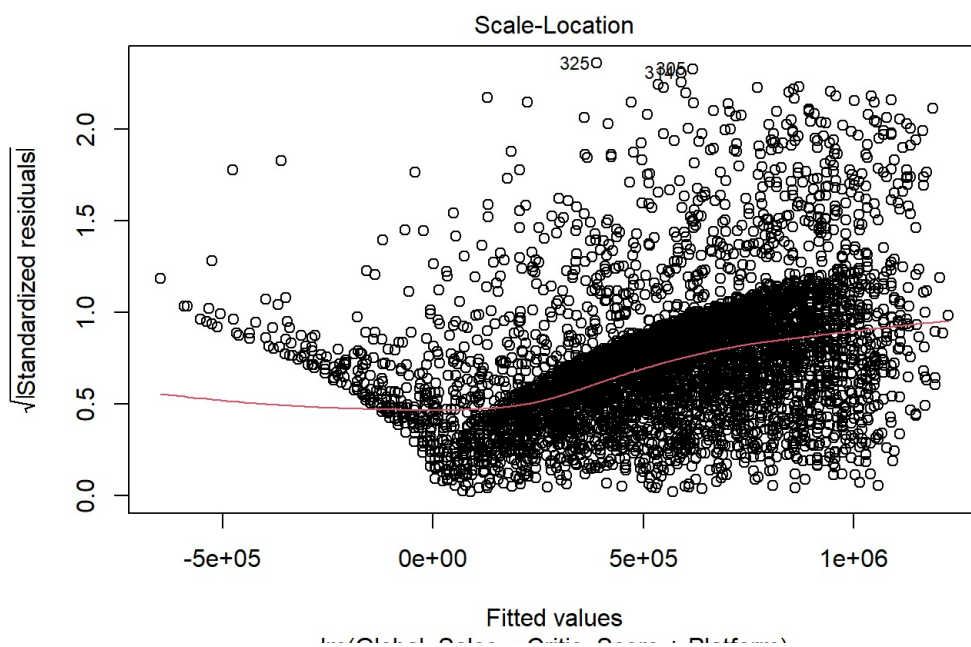
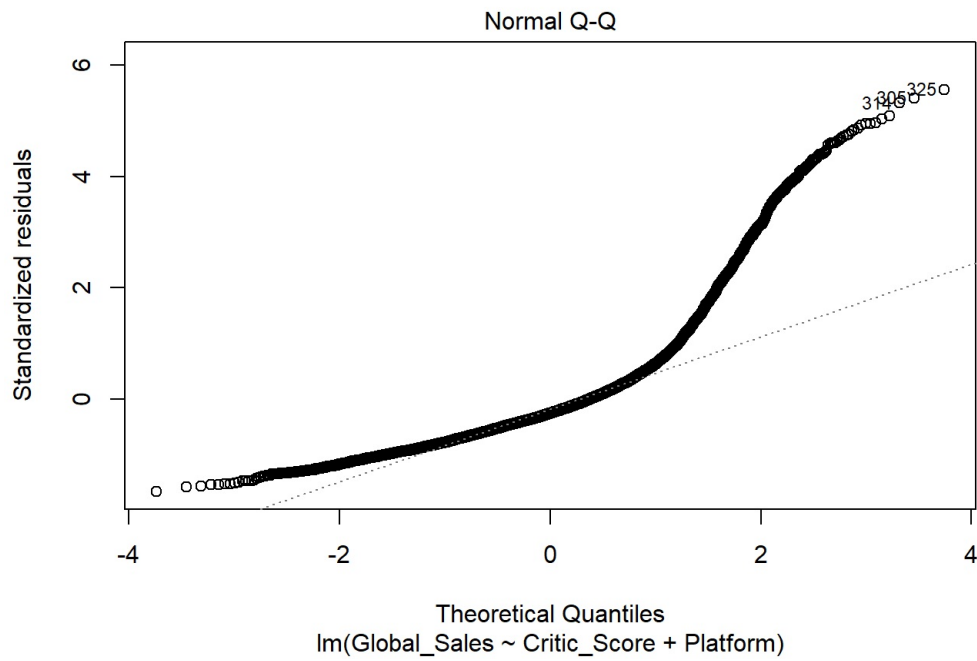
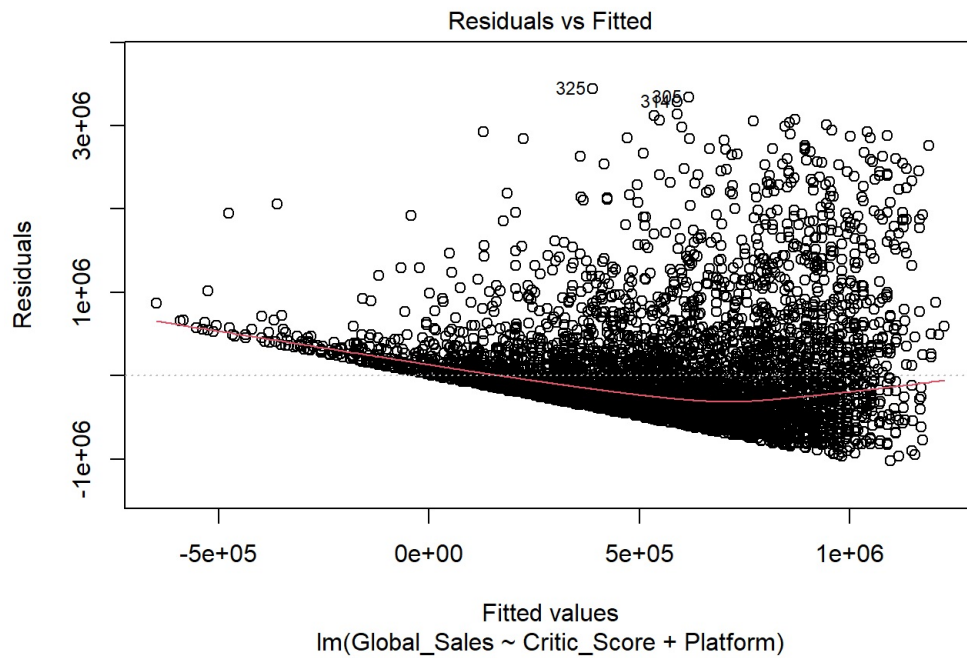
```
lm2 <- lm(Global_Sales~Critic_Score+Platform, data=train)
summary(lm2)
```

```
##
## Call:
## lm(formula = Global_Sales ~ Critic_Score + Platform, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1026450 -377974 -156053  164029  3440888
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -795845.2    70456.7  -11.296 < 2e-16 ***
## Critic Score   19225.3     631.3   30.454 < 2e-16 ***
## PlatformDC    -532201.1   181072.7  -2.939 0.003305 **
## PlatformDS      50663.6    64538.3   0.785 0.432479
## PlatformGBA    -91686.8    71428.1  -1.284 0.199328
## PlatformGC    -183501.6    66835.7  -2.746 0.006061 **
## PlatformPC    -421095.3    62019.8  -6.790 1.24e-11 ***
## PlatformPS     155797.7    80237.4   1.942 0.052224 .
## PlatformPS2    113586.7    59569.1   1.907 0.056598 .
## PlatformPS3    181242.9    61265.4   2.958 0.003107 **
## PlatformPS4     67723.5    71372.7   0.949 0.342729
## PlatformPSP    -92447.1    65871.6  -1.403 0.160541
## PlatformPSV   -314549.6    82993.2  -3.790 0.000152 ***
## PlatformWii    195963.0    64687.9   3.029 0.002462 **
## PlatformWiiU   -86323.8    93194.4  -0.926 0.354343
## PlatformX360   174174.0    60717.2   2.869 0.004139 **
## PlatformXB    -216947.6    62969.6  -3.445 0.000575 ***
## PlatformXOne   74272.8     78427.9   0.947 0.343671
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 619500 on 5441 degrees of freedom
## Multiple R-squared:  0.1922, Adjusted R-squared:  0.1897
## F-statistic: 76.15 on 17 and 5441 DF, p-value: < 2.2e-16
```

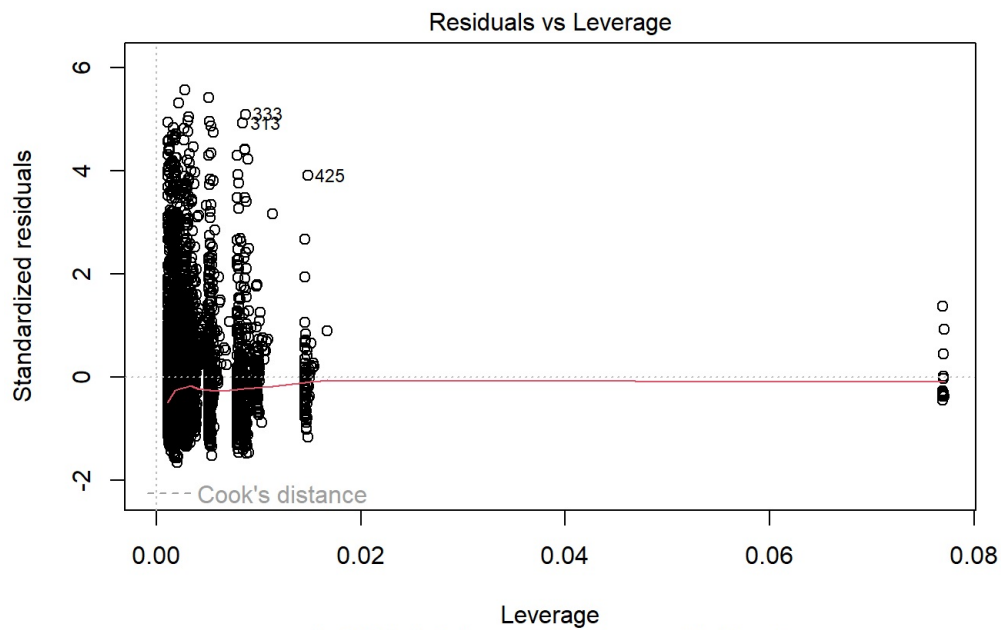

Overall, the model saw a slight improvement. Most notably, the Adjusted R-squared value increased to about .19. Its still not a good model, but it has improved. The residual error also reduced by tens of thousands. The range of min and max didn't improve much. Our F statistic went down, but the p value remained low.

Now the residual will be plotted.

```
plot(lm2)
```



lm(Global_Sales ~ Critic_Score + Platform)



Unfortunately, all of the new graphs

lm(Global_Sales ~ Critic_Score + Platform)

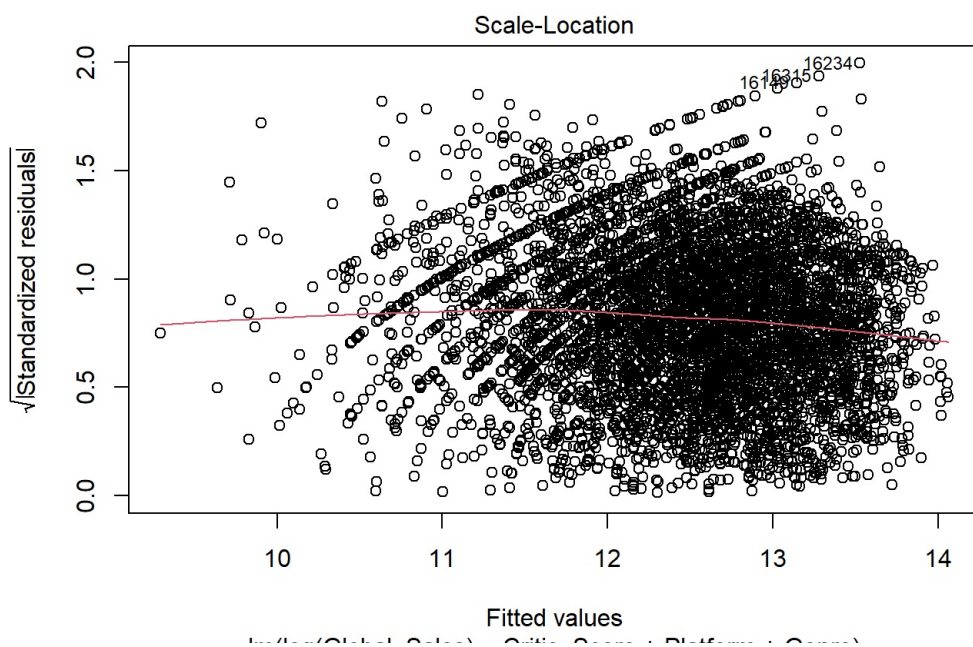
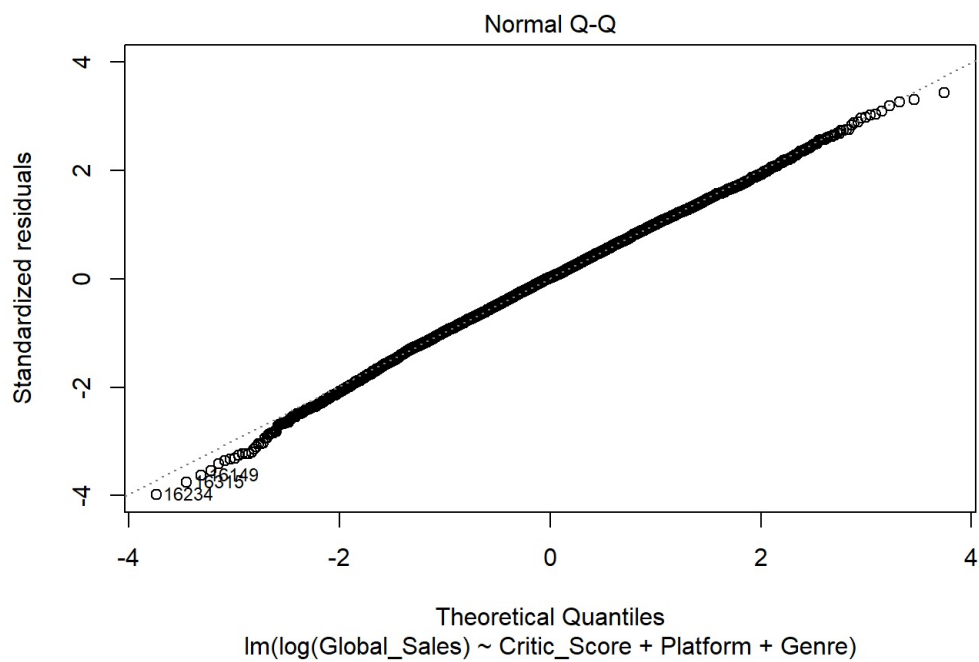
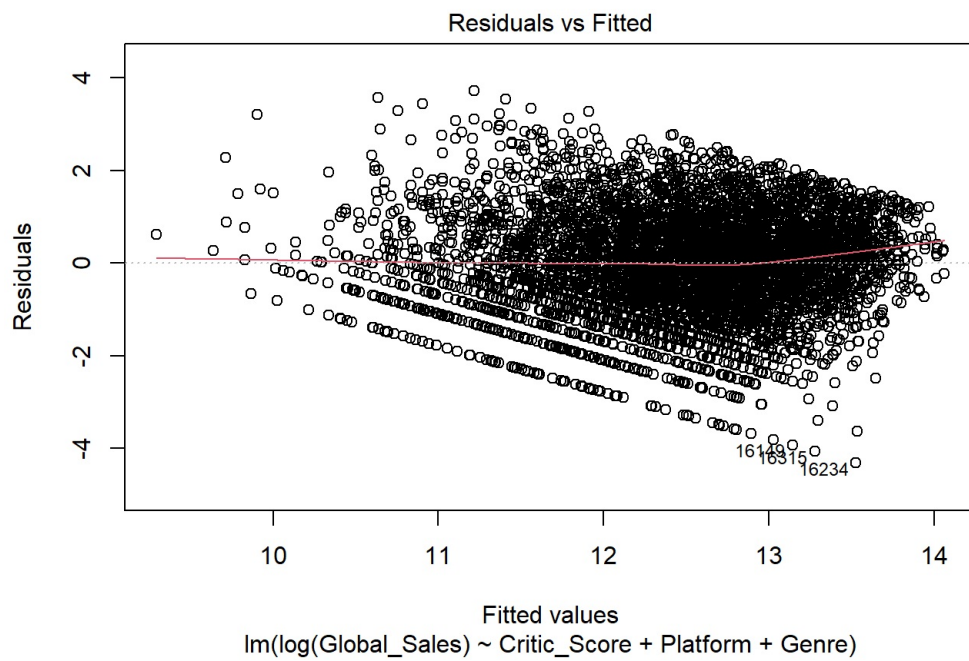
resemble the old graphs and have the same issues as described above. Despite the small improvement, the model is still bad, as suspected.

g:

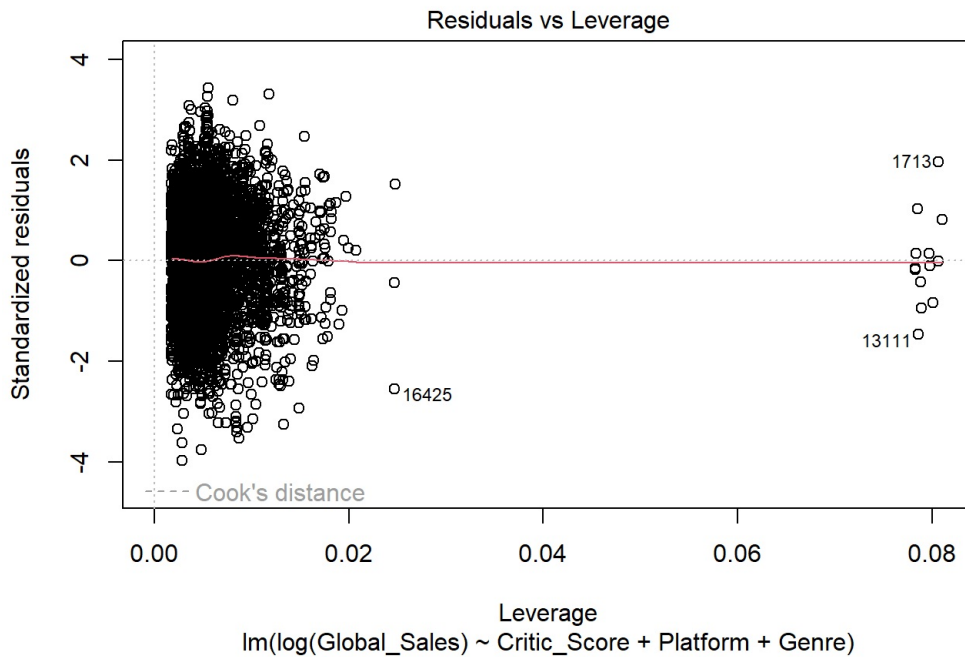
```
lm3 <- lm(log(Global_Sales)~Critic_Score+Platform+Genre, data=train)
summary(lm3)
```

```
##
## Call:
## lm(formula = log(Global_Sales) ~ Critic_Score + Platform + Genre,
##     data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.3166 -0.7136  0.0200  0.7381  3.7162
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   10.058361   0.124887  80.540 < 2e-16 ***
## Critic_Score    0.038572   0.001123  34.354 < 2e-16 ***
## PlatformDC     -0.927697   0.317829  -2.919 0.003528 **
## PlatformDS     -0.057235   0.113643  -0.504 0.614538
## PlatformGBA    -0.323553   0.125631  -2.575 0.010038 *
## PlatformGC     -0.492399   0.117854  -4.178 2.99e-05 ***
## PlatformPC     -1.563879   0.110103 -14.204 < 2e-16 ***
## PlatformPS      0.212093   0.140918   1.505 0.132363
## PlatformPS2     0.144849   0.105071   1.379 0.168082
## PlatformPS3     0.313780   0.108018   2.905 0.003689 **
## PlatformPS4    -0.336518   0.125511  -2.681 0.007358 **
## PlatformPSP     -0.262939   0.115896  -2.269 0.023323 *
## PlatformPSV    -0.698278   0.145694  -4.793 1.69e-06 ***
## PlatformWii      0.322858   0.113888   2.835 0.004601 **
## PlatformWiiU    -0.268102   0.163642  -1.638 0.101408
## PlatformX360     0.239864   0.107176   2.238 0.025258 *
## PlatformXB     -0.583072   0.111546  -5.227 1.79e-07 ***
## PlatformXOne    -0.251256   0.138195  -1.818 0.069100 .
## GenreAdventure  -0.589257   0.081135  -7.263 4.33e-13 ***
## GenreFighting   -0.070778   0.069312  -1.021 0.307228
## GenreMisc        0.129744   0.069630   1.863 0.062470 .
## GenrePlatform    0.008523   0.068722   0.124 0.901304
## GenrePuzzle     -0.551773   0.117243  -4.706 2.59e-06 ***
## GenreRacing     -0.158295   0.059391  -2.665 0.007714 **
## GenreRole-Playing -0.211605   0.055544  -3.810 0.000141 ***
## GenreShooter     0.018727   0.051860   0.361 0.718027
## GenreSimulation  0.019945   0.077922   0.256 0.797994
## GenreSports     -0.046738   0.051004  -0.916 0.359514
## GenreStrategy   -0.593321   0.081109  -7.315 2.95e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.086 on 5430 degrees of freedom
## Multiple R-squared:  0.3126, Adjusted R-squared:  0.309
## F-statistic: 88.18 on 28 and 5430 DF, p-value: < 2.2e-16
```

```
plot(lm3)
```



lm(log(Global_Sales) ~ Critic_Score + Platform + Genre)



h: The third model is by far the best one. It added genre as a factor, and used the log of the sales.

The variation from the median is still large (or at least, large for a log function anyways). However, the residual error decreased a bit while the Adjusted R-squared jumped up to .30. Still not great, but its an improvement over .11 and .19 of the first two models. Its unfortunate that our F statistic keeps getting lower, but the p value remains low as well.

As far as residuals go, the third also shows great improvements.

Residuals vs. fitted: The line becomes almost perfectly horizontal (except for the end), and has a roughly symmetrical distribution. This indicates the linear relationship is a good match.

Normal Q-Q: The residuals are plotted pretty well against the line. This means they are distributed fairly normally.

Scale location: This one is only so-so. The line in the middle is decently horizontal, and distribution follows a roughly ovular distribution around the line. It's by no means great, but a patters can sort of be recognized. This would mean the residual are spread out decently among the range of predictors.

Residuals vs Leverage: Unfortunately, all the data seems to be split by a line. This means several factors have too much sway over the data.

Overall, the third model shows many improvements over the other 2.

i. Now lm3 will be used to compare predictions with th e test data.

```
pred1 <- predict(lm1, newdata = test)
corr1 <- cor(pred1, test$Global_Sales)
mse1 <- mean((pred1-test$Global_Sales)^2)
rmse1 <- sqrt(mse1)
```

```
print("lm1")
```

```
## [1] "lm1"
```

```
print(paste("correlation: ", corr1))
```

```
## [1] "correlation: 0.270310200602469"
```

```
print(paste("mse: ", mse1))
```

```
## [1] "mse: 5198864194363.64"
```

```
print(paste("rmse: ", rmse1))
```

```
## [1] "rmse: 2280101.79473717"
```

```
pred2 <- predict(lm2, newdata = test)
corr2 <- cor(pred2, test$Global_Sales)
mse2 <- mean((pred2-test$Global_Sales)^2)
rmse2 <- sqrt(mse2)

print("lm2")
```

```
## [1] "lm2"
```

```
print(paste("correlation: ", corr2))
```

```
## [1] "correlation: 0.335563801521632"
```

```
print(paste("mse: ", mse2))
```

```
## [1] "mse: 5039282962851.88"
```

```
print(paste("rmse: ", rmse2))
```

```
## [1] "rmse: 2244834.72951839"
```

```
pred3 <- predict(lm3, newdata = test)
corr3 <- cor(pred3, test$Global_Sales)
mse3 <- mean((pred3-test$Global_Sales)^2)
rmse3 <- sqrt(mse3)

print("lm3")
```

```
## [1] "lm3"
```

```
print(paste("correlation: ", corr3))
```

```
## [1] "correlation: 0.322831823853248"
```

```
print(paste("mse: ", mse3))
```

```
## [1] "mse: 6108345435624.09"
```

```
print(paste("rmse: ", rmse3))
```

```
## [1] "rmse: 2471506.71365143"
```

The results were a bit unexpected.

The second model had the highest correlation with the test data, with the third model not too far behind it. Though .33 still isn't great.

The third model had a significantly worse mse than the other two.

The rmse was roughly the same, but worse in the third model.

The second model must be the best mix of increasing complexity while maintaining accuracy with the randomness of the data. It had two factors, which was in between the 1 and 3 factors of the others. It also didn't take the log of the values, which seems to have made the errors grow larger.

My guess is overall the data itself doesn't have enough information to successfully determine video game sales. It makes sense because things such as the actual price of the game, if games of a similar type release at the same time, if sufficient advertising was done, if the game can only be sold on certain consoles, etc. will affect sales. Considering all of that, I think its reasonable to assume reviews will only influence about 30% of a games sales.